

AI Unleashed: Autonomous Attacks and Self-Healing Defenses

Bimal Kandel

BSc (Hons) Cybersecurity
Herald College Kathmandu
Kathmandu, Nepal
np03cy4a240071@heraldcollege.edu.np

Shaurab Thapa

BSc (Hons) Cybersecurity
Herald College Kathmandu
Kathmandu, Nepal
np03cy4a240045@heraldcollege.edu.np

Abstract: The fast progress of artificial intelligence (AI) creates a new type of cybersecurity threat which consists of self-governing AI systems that launch continuous network assaults because they operate at the speed of machines. The research paper investigates the methods and systems through which autonomous artificial intelligence systems perform intelligent network attacks which include reinforcement learning (RL)-driven exploitation and generative adversarial network (GAN)-based evasion and self-optimizing polymorphic malware. The study investigates dependable defense systems which possess self-repair capabilities. These systems use deep learning techniques to identify unusual patterns while they share threat data and manage security incidents through automatic processes. The current defenses face fundamental weaknesses which we discover through advanced research. Our security framework consists of multiple protection layers which adapt to new threats. It enables real-time detection and threat mitigation and system recovery from AI-based attacks. The study investigates two main areas: adversarial robustness methods and explainable AI (XAI) as a defense tool and the connection between artificial intelligence and zero-trust security systems. The research demonstrates that security systems need to switch to AI-powered protective measures which must adapt to emerging threats.

Keywords: autonomous cyberattacks, AI-driven intrusion, adversarial machine learning, self-healing networks, anomaly detection, federated learning, zero-trust architecture, cyber resilience

I. INTRODUCTION

Current digital environments extend their reach beyond existing enterprise network frameworks. The combination of cloud-native designs and Internet of Things (IoT) and edge computing and 5G networks creates an extremely large security vulnerability area. The existing signature-based intrusion detection systems together with static firewall security mechanisms demonstrate their fundamental weakness against attackers who utilize zero-day exploits and advanced persistent threats (APTs) and polymorphic malware which evolves at a faster rate than signature databases can refresh [1].

The present security landscape shows major changes because of AI which now serves as a combat tool. Present-day autonomous attack systems possess the ability to explore computer networks in order to find weak security areas which they can exploit while developing their attack methods to match the current security countermeasures that protect the system from human-based monitoring. Mohamed's comprehensive review [1] of AI and machine learning in cybersecurity demonstrates that adversarial AI has taken control of the entire cyber kill chain which begins with reconnaissance and initial access and ends with data exfiltration through lateral movement and privilege escalation.

The defensive community has not stood still. Security systems now use deep learning models together with federated threat intelligence platforms and automatic security orchestration systems. The existing security situation maintains an unbalanced state because attackers only need one successful attempt while defenders must achieve flawless performance on each occasion. The solution to this problem requires security systems which combine active defense mechanisms with intelligent defense capabilities that predict attacker behavior through their real-time learning abilities which enable them to process security incidents and automatically restore system functions.

The paper makes its contributions through four distinct research activities which include: (1) performing a systematic study of AI-based attack methods; (2) introducing a self-healing defense system which uses current academic research as its foundation; (3) testing current security systems against authentic performance standards; and (4) investigating future research possibilities which include explainable artificial intelligence and zero-trust security systems and quantum-resistant cryptographic methods. The remainder of the paper is organized as follows: Section II reviews related work; Section III analyzes autonomous

attack mechanisms; Section IV presents self-healing defense architectures; Section V discusses evaluation and findings; Section VI explores additional defense dimensions; and Section VII concludes.

II. RELATED WORK

Research at the intersection of AI and cybersecurity has increased exponentially during the last five years. The initial supervised learning methods used traditional classifiers to test intrusion detection datasets which included KDD Cup 99 and NSL-KDD through their decision trees and random forests and support vector machines which achieved high accuracy results for established attack types. The models demonstrated their weakness when they encountered new attack types which had not been included in their training materials that became more crucial as attackers learned how to bypass machine learning detection systems.

The recent survey of this field conducted by Mohamed [1] provides the most detailed examination which includes deep learning methods for malware classification and behavioral analytics techniques used to detect insider threats and federated learning approaches that sustain privacy during joint defensive work and adversarial robustness methods. The review explicitly identifies the need for adaptive adversarial defense systems and federated learning frameworks capable of global threat mitigation—the central themes of the present paper.

Awadallah et al. [2] conducted simultaneous research which showed that autonomous threats present a danger that extends to both traditional networks and new environments which include the metaverse and extended reality platforms. This finding demonstrates that complete defense systems need to operate independently from particular network designs. Samtani et al. [4] developed explainable artificial intelligence for cyber threat intelligence systems (XAI-CTI) to show that defense systems need to provide transparent AI decision-making processes which analysts require to establish trust and meet regulatory standards.

On the offensive side, reinforcement learning-based penetration testing agents have been studied by multiple groups, with Riedl [6] providing a thorough survey of RL approaches for autonomous cyber operations. These agents demonstrate that AI-driven attacks are not theoretical constructs but practical tools whose capabilities already exceed those of many deployed defenses. This body of work collectively motivates both the threat analysis and defense proposals in the sections that follow.

III. AI-POWERED ATTACK STRATEGIES

A. Reinforcement Learning-Driven Exploitation

Reinforcement learning (RL) lets autonomous agents find the best attack paths by repeatedly testing their abilities in different target environments. The RL agent uses the reward function to test various network exploits while it learns from the results of its tests. This system runs at human red-team proficiency because it can perform thousands of its testing processes every second.

RL-based attack agents show unpredicted behavior through their ability to discover multi-hop pivot pathways, their capacity to time attacks according to periods with reduced security checks, and their capability to merge various minor security weaknesses into a major system breach. The agent shows adaptive capacity because it reroutes after defensive measures succeed in stopping one pathway which enables autonomous attackers to operate differently from fixed exploit scripts [3, 6].

TABLE I. AI ATTACK TECHNIQUES AND CORRESPONDING DEFENSES

Attack Type	AI Mechanism	Key Defense
RL Exploitation	Policy gradient	Anomaly IDS
GAN Evasion	Adversarial gen.	Adv. training
Polymorphic Malware	Self-mutation	Behavioral ML
LLM Phishing	Text generation	NLP filters
Botnet Swarm	Federated attack	FL defense

B. GAN-Based Evasion and Adversarial Samples

Generative adversarial networks (GANs) have been repurposed as offensive tools to craft malicious payloads that evade ML-based classifiers. The generator produces adversarial samples network packets, executable binaries, or document payloads, while the discriminator evaluates their statistical similarity to benign traffic. Through iterative training, the generator learns to produce inputs that fool both the discriminator and real-world intrusion detection systems [5].

Mohamed's review [1] identifies GAN-driven malware obfuscation as one of the most technically challenging offensive developments, noting that generated samples can reduce the detection accuracy of state-of-the-art classifiers by 30–50%. Since generated samples are novel by construction, they evade signature matching entirely, and their statistical indistinguishability from benign traffic undermines threshold-based anomaly detectors.

C. Self-Optimizing and Polymorphic Malware

Beyond initial intrusion, AI enables malware to evolve post-deployment. Self-optimizing malware embeds lightweight inference engines within its payload to assess the local environment, identify active monitoring tools, and mutate code segments to evade detection. Unlike classical polymorphic engines that simply re-encrypt payloads with changing keys, AI-driven variants optimize their mutations against learned models of the target's detection stack, producing variants specifically tailored to bypass the defenses actually deployed on the victim host.

Across a botnet, such malware coordinates via distributed channels, sharing evasion strategies learned on individual compromised nodes. This collective intelligence transforms a collection of infected machines into a collaborating adversarial swarm, each contributing to a shared evasion policy that defeats defenses trained on older samples [1].

D. Autonomous Social Engineering

Large language models (LLMs) have introduced a natural language dimension to autonomous attacks. AI-powered phishing agents scrape public social media and professional profiles to construct highly personalized spear-phishing emails, voice messages, and chat interactions at industrial scale, without human authorship. The contextual accuracy of LLM-generated lures dramatically increases engagement rates compared to generic phishing, and the automation makes volume-based defenses such as rate limiting insufficient. This attack vector bypasses technical controls entirely, targeting the human element of the security perimeter.

IV. SELF-HEALING DEFENSE MECHANISMS

A. Deep Learning Anomaly Detection

Effective defense against AI-driven attacks requires continuous detection through real-time behavior analysis instead of traditional signature-based detection methods. Long Short-Term Memory (LSTM) networks together with autoencoder architectures demonstrate their ability to track network traffic flow while identifying security breaches that occur without predefined attack patterns [1].

The autoencoder model learns to duplicate standard network operations through its reconstruction process which aims for minimal output deviation from actual network activity. The system detects unusual behavior when it receives input that results in substantial reconstruction errors which include detection of port scanning activities and lateral network movement and hidden data transfer operations. The system calculates reconstruction error through the application of Mean Squared Error (MSE) loss function. The anomaly score for an input sample x with n features and reconstructed output $\hat{x}$, the anomaly score is defined as:

$$L(x) = (1/n) \times \sum (x_i - \hat{x}_i)^2 \quad [1]$$

where x_i represents the i -th feature of the original traffic vector while $\hat{x}_i$ shows the corresponding reconstructed value produced by the autoencoder decoder. The system classifies a sample as an intrusion when $L(x)$ exceeds a threshold of θ which the system learned from validation traffic data. This method does not require any attack data which enables it to identify zero-day exploits together with new GAN-based evasion methods that would bypass signature detection systems [7].

B. Federated Threat Intelligence

Federated learning (FL) enables organizations to train a shared detection model collaboratively without exchanging raw network data [8]. Each participant trains locally and contributes only gradient updates to a central aggregator, preserving the confidentiality of sensitive operational traffic. In the context of AI-powered attacks, federated models are particularly valuable because attack campaigns rarely target isolated organizations. An RL attack agent that succeeds against one enterprise will probe others using similar techniques; a federation trained on hundreds of organizations' experiences will recognize these patterns far earlier than any single-organization model.

C. Automated Incident Response

Self-healing networks couple anomaly detection with automated response orchestration. Upon detection, the response pipeline executes tiered playbooks: micro-segmenting affected network zones, revoking compromised credentials, rolling back configurations to verified snapshots, and generating human-readable incident reports—all within seconds of detection. An RL-based response agent governs this pipeline, learning from past incidents to optimize containment decisions, minimize service disruption, and reduce false-positive isolation events. This architecture reduces mean time to respond (MTTR) from hours to seconds [1].

D. Adversarial Model Hardening

Since AI-powered attacks specifically target ML-based detectors, hardening models against adversarial manipulation is essential. Adversarial training incorporates GAN-generated adversarial samples into the training corpus, improving robustness against evasion at inference time. Certified defenses using randomized smoothing provide theoretical accuracy guarantees within bounded perturbation radii, offering principled protection that complements empirical adversarial training [1]. Ensemble methods that combine multiple independently trained detectors further reduce the probability that a single adversarial perturbation defeats all component models simultaneously.

V. EVALUATION AND DISCUSSION

The current assessments of AI-based offensive and defensive systems show both strengths and weaknesses. RL penetration testing agents have achieved hacking success rates above 80% against unpatched enterprise network systems in their offensive operations [6]. GAN-based evasion studies have shown that classifier detection accuracy decreases by 30 to 50 percent when testing models trained on regular data without adversarial testing [1].

Table II shows how four different IDS system types performed detection tests on two benchmark datasets CICIDS2017 and NSL-KDD. LSTM-based anomaly detectors achieve the highest detection rate (96%) while maintaining the lowest false positive rate (1.8%), confirming the advantage of deep temporal modeling over rule-based approaches. Federated LSTM enables multiple organizations to deploy their systems while maintaining identical accuracy results. Automated response systems enable organizations to decrease mean time to respond (MTTR) by two to three orders of magnitude when they use these systems instead of human-operated SOC processes.

TABLE II. IDS PERFORMANCE COMPARISON (CICIDS2017 / NSL-KDD)

IDS System	Det. Rate	FP Rate	MTTR
Signature-Based	72%	8.3%	~4 hrs.
Rule-Based IDS	81%	5.7%	~2 hrs.
LSTM Anomaly Det.	96%	1.8%	~12 sec
Federated LSTM	95%	2.1%	~15 sec

The actual implementation of machine learning systems still faces its main challenges despite the achieved breakthroughs. The first point demonstrates that benchmark results fail to predict actual performance in production systems because real-world traffic patterns experience ongoing changes while static models need continuous retraining to maintain their accuracy. The second point establishes that deep learning systems require more efficient processing methods because their current latency issues create problems for high-volume operations which drives research efforts to develop both lightweight processing solutions and hardware-based inference systems. The current situation of the adversarial arms race gives attackers an advantage because attackers need to find and exploit only one vulnerability while defenders must protect all potential entry points. The existing technological solutions fail to eliminate this power imbalance because artificial intelligence security systems depend on human skills and partnerships between different industry sectors for their effective operation.




Fig. 1. IDS Detection Rate vs. False Positive Rate Comparison.

VI. ADVANCED DEFENSE DIMENSIONS

A. Zero-Trust Architecture Integration

Zero-trust architecture (ZTA) applies its core principle which denies any user device or network segment access and trust regardless of their current physical location. The system requires all users to undergo authentication and authorization processes which maintain security through ongoing assessment of their established behavioral patterns. The combination of AI-powered anomaly detection with ZTA system creates a security defense mechanism which maintains its effectiveness even when an attacker possesses legitimate access rights because all security violations will lead to increased authentication requirements or session termination, which prevents further movement across systems. Mohamed [1] identifies ZTA as a critical complement to AI-based detection, particularly in hybrid and multi-cloud environments where perimeter-based security lacks structural viability.

ZTA requires organizations to implement three essential components which include precise identity control and constant security monitoring of all devices and automated security policy implementation. ML classifiers enable differentiation between valid credential usage and credential-stuffing attacks through their analysis of typing patterns and session durations and access behavior consistency which functions as invisible user authentication.

B. Explainable AI for Transparent Defense

The deep learning security systems face ongoing criticism because their internal workings remain hidden, which creates difficulties for analysts who need to understand LSTM-triggered anomalies and RL-agent network segmentations. The system's hidden operations create two problems which undermine trustworthiness and make it harder to meet regulatory requirements while risking misidentifying authentic security threats as nonthreatening events. The Explainable AI (XAI) techniques which include attention mechanisms and SHAP (SHapley Additive exPlanations) values and counterfactual explanations, present the precise features and behavior patterns which resulted in each detection decision [4].

The research by Samtani et al. [4] demonstrates that XAI-enabled cyber threat intelligence systems help analysts better respond to threats while decreasing their alert fatigue when compared to black-box detectors. The need for explainability extends beyond mere usability in automated defense systems because it serves as an ethical requirement which requires organizations to explain their decision-making process for actions that can harm legitimate business operations.

C. Quantum-Resilient Cryptographic Foundations

The combination of artificial intelligence and quantum computing technology creates a new permanent danger which will endanger existing systems of cryptographic security. The

quantum algorithms which include Shor's algorithm can perform integer factorization and discrete logarithm problem-solving at polynomial speed, which makes RSA and ECC encryption systems vulnerable to potential future quantum computing power. Mohamed [1] identifies quantum-AI convergence as one of the most significant frontier challenges in cybersecurity, requiring organizations to begin cryptographic agility planning now, well before practical quantum computers reach adversarial hands.

NIST currently finalizes post-quantum cryptographic standards which eliminate algorithms that quantum systems can break through their introduction of lattice-based, hash-based, and code-based methods that remain impossible to solve by both classic and quantum attackers [9]. AI can accelerate the migration process by automatically inventorying cryptographic assets across large enterprise environments, prioritizing migration by risk exposure, and detecting the use of deprecated algorithms in real time.

D. Threat Hunting and Proactive Defense

The standard operation of security operations centers (SOCs) requires security personnel to respond to security threats which emerge from their monitoring systems. AI-enabled threat hunting inverts this posture because automated hunting agents use network telemetry to identify early indicators of compromise (IoCs) and known adversary techniques (TTPs) instead of waiting for detectors to activate. Hunting agents use temporal and network segment analysis to detect APTs through their reconnaissance stages because they can detect weak signals that exist below detection limits [1, 10].

Natural language processing (NLP) models extend their functionality by processing open-source threat intelligence feeds and dark web forums and vulnerability disclosures to detect upcoming attack tools and techniques which will be used against the organization. The organization has changed its intelligence approach from perimeter defense methods to continuous methods which model adversary behavior.

VII. CONCLUSION

The paper investigates how artificial intelligence functions as a dual system within network security through its capacity to launch automatic learning attacks which continuously alter their strategies and its ability to create self-repairing networks that maintain operational strength.

AI-powered attack strategies RL-driven exploitation, GAN-based evasion, polymorphic malware, and autonomous social engineering represent a qualitative escalation in the threat landscape that demands equally capable, equally adaptive defenses. The research team developed a complete AI-based defense framework which combines six security components. The framework overcomes the main drawbacks which exist in signature-based detection systems and static detection systems by offering dynamic security protection which adapts to new threats.

Researchers and practitioners should focus on three main research areas. The first requirement states that automated defense systems must show their actions through clear and verifiable explanations which build analyst confidence while fulfilling regulatory obligations. The second point requires organizations to start their post-quantum cryptographic transition through AI-based asset evaluation and risk assessment for their assets. The third point establishes that international security operations need to develop shared global frameworks which allow countries to share intelligence about threats, while applying privacy protection methods that guarantee faster transmission of knowledge from defenders to attackers. The development of cybersecurity needs to focus on three main elements which include autonomous systems, cooperative

methods and explainable systems through persistent funding of AI-driven security research.

ACKNOWLEDGMENT

We extend our gratitude to Herald College Kathmandu and the Faculty of BSc (Hons) Cybersecurity, University of Wolverhampton, for their ongoing academic support to our research work. The open-access research community deserves our recognition because their published research enabled us to conduct this review.

REFERENCES

- [1] N. Mohamed, "Artificial intelligence and machine learning in cybersecurity: a deep dive into state-of-the-art techniques and future paradigms," *Knowledge and Information Systems*, vol. 67, pp. 6969–7055, Apr. 2025.
<https://link.springer.com/article/10.1007/s10115-025-02429-y>
- [2] A. Awadallah et al., "Artificial intelligence-based cybersecurity for the metaverse: research challenges and opportunities," *IEEE Communications Surveys and Tutorials*, 2024.
<https://www.researchgate.net/publication/383076146>
- [3] G. Palmer, C. Parry, D. J. B. Harrold, and C. Willis, "Deep reinforcement learning for autonomous cyber operations: a survey," *arXiv preprint arXiv:2310.07745*, 2024.
<https://arxiv.org/abs/2310.07745>
- [4] S. Samtani, H. Chen, M. Kantarcioglu, and B. Thuraisingham, "Explainable artificial intelligence for cyber threat intelligence (XAI-CTI)," *IEEE Transactions on Dependable and Secure Computing*, vol. 19, no. 4, pp. 2149–2150, 2022.
<https://ieeexplore.ieee.org/document/9821359>
- [5] I. J. Goodfellow et al., "Generative Adversarial Nets," in *Advances in Neural Information Processing Systems*, vol. 27, 2014.
<https://arxiv.org/pdf/1406.2661>
- [6] G. Palmer et al., "Deep reinforcement learning for autonomous cyber defence: a survey," *arXiv preprint arXiv:2310.07745v2*, Sep. 2024.
<https://arxiv.org/pdf/2310.07745v2>
- [7] H. Sedjelmaci, S. M. Senouci, and T. Taleb, "An accurate and efficient collaborative intrusion detection framework to secure vehicular networks," *IEEE Transactions on Intelligent Transportation Systems*, vol. 18, no. 5, pp. 1074–1085, 2017.
<https://www.sciencedirect.com/science/article/abs/pii/S004579061500066X>
- [8] B. McMahan et al., "Communication-efficient learning of deep networks from decentralized data," in *Proc. 20th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, vol. 54, 2017, pp. 1273–1282.
<https://arxiv.org/pdf/1602.05629>
- [9] NIST, "Post-quantum cryptography standardization," National Institute of Standards and Technology, *NIST IR 8413*, 2022. [Online]. Available: <https://csrc.nist.gov/projects/post-quantum-cryptography>
- [10] A. Sarker et al., "On autonomous agents in a cyber defence environment," *arXiv preprint arXiv:2309.07388*, 2023.
<https://arxiv.org/pdf/2309.07388>